

Assignment 3 – Report

Group - 13

Ahmed Abdelmaksoud

Hassan Ahmed

Dina Abdelhady

Ahmed Abdelsamad

1. Event Hubs Analytics

1. Create a resource group

The screenshot shows the Microsoft Azure portal interface. The top navigation bar includes the Microsoft Azure logo, a search bar, and user information. The left sidebar contains navigation links for Overview, Activity log, Access control (IAM), Tags, Resource visualizer, Events, Settings, Deployments, Security, Policies, Properties, Locks, Cost Management, Cost analysis, Cost alerts (preview), Budgets, and Advisor recommendations. The main content area displays the 'ELG5166_Resources' resource group. The 'Essentials' section shows the Subscription (Move) as 'Azure for Students', Subscription ID as '02053776-f598-4f39-b8f4-cbe11900e07e', and Location as 'East US'. The 'Resources' section shows a list of resources with columns for Name, Type, and Location. The resources listed are: assignment3 (Data Lake Storage Gen1, East US 2), assignment3 (Event Hubs Namespace, East US), assignment3 (Data Lake Analytics, East US 2), bikedatajob (Stream Analytics job, East US), and bikedatareceiver123 (Storage account, East US).

2. Create event hub name space

The screenshot shows the Microsoft Azure portal interface for the 'assignment3' Event Hubs Namespace. The top navigation bar includes the Microsoft Azure logo, a search bar, and user information. The left sidebar contains navigation links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Events, Settings, Shared access policies, Scale, Geo-Recovery, Encryption, Properties, and Locks. The main content area displays the 'assignment3' Event Hubs Namespace. The 'Essentials' section shows the Resource group (Move) as 'ELG5166_Resources', Status as 'Active', Location as 'East US', Subscription (Move) as 'Azure for Students', Subscription ID as '02053776-f598-4f39-b8f4-cbe11900e07e', and Host name as 'assignment3.servicebus.windows.net'. The 'Tags (Edit)' section shows a link to 'Click here to add tags'. The 'Summary' section shows the namespace contents, Kafka surface, and zone redundancy. The namespace contents are: 1 EVENT HUB. The Kafka surface is: NOT SUPPORTED. The zone redundancy is: ENABLED. The 'Show data for the last:' section shows a dropdown menu with options: 1 hour, 6 hours, 12 hours, 1 day, 7 days, and 30 days. The 'Requests' and 'Throughput' sections are visible at the bottom.

3. Create event hub

Microsoft Azure | Search resources, services, and docs (G+/I)

Home > assignment3 > **bike (assignment3/bike)** Event Hubs Instance

Search (Ctrl+/) < + Consumer group Delete Refresh

Overview

Access control (IAM)

Diagnose and solve problems

Settings

Shared access policies

Properties

Locks

Entities

Consumer groups

Features

Capture

Process data

Automation

Tasks (preview)

Essentials

Resource group (Move) : ELG5166_Resources

Location : East US

Subscription (Move) : Azure for Students

Subscription ID : 02053776-f598-4f39-b8f4-cbe1190e07e

Partition Count : 2

Status : Active

Namespace : assignment3

Created : Friday, November 26, 2021, 18:36:33 GMT+2

Updated : Saturday, November 27, 2021, 15:53:16 GMT+2

Message Retention : 1 day

Capture events
Use Capture to save your events to persistent storage.

Process data
Produce insights with Azure's data processing services.

Connect
Authenticate with connection strings and SAS policies.

Checkpoint
Create consumer groups to checkpoint your events.

JSON View

4. Copy connection string key.

Microsoft Azure | Search resources, services, and docs (G+/I)

Home > assignment3 > **assignment3 | Shared access policies** Event Hubs Namespace

Search (Ctrl+/) < + Add

Search to filter items...

Policy

RootManageSharedAccessKey

Claims

Manage, Send, Listen

SAS Policy: RootManageShare...

Save Discard Delete ...

Manage

Send

Listen

Primary key
Mk23UP1a+mT4EYwZ7ea3J0w31LZsws+hhscL378g4=

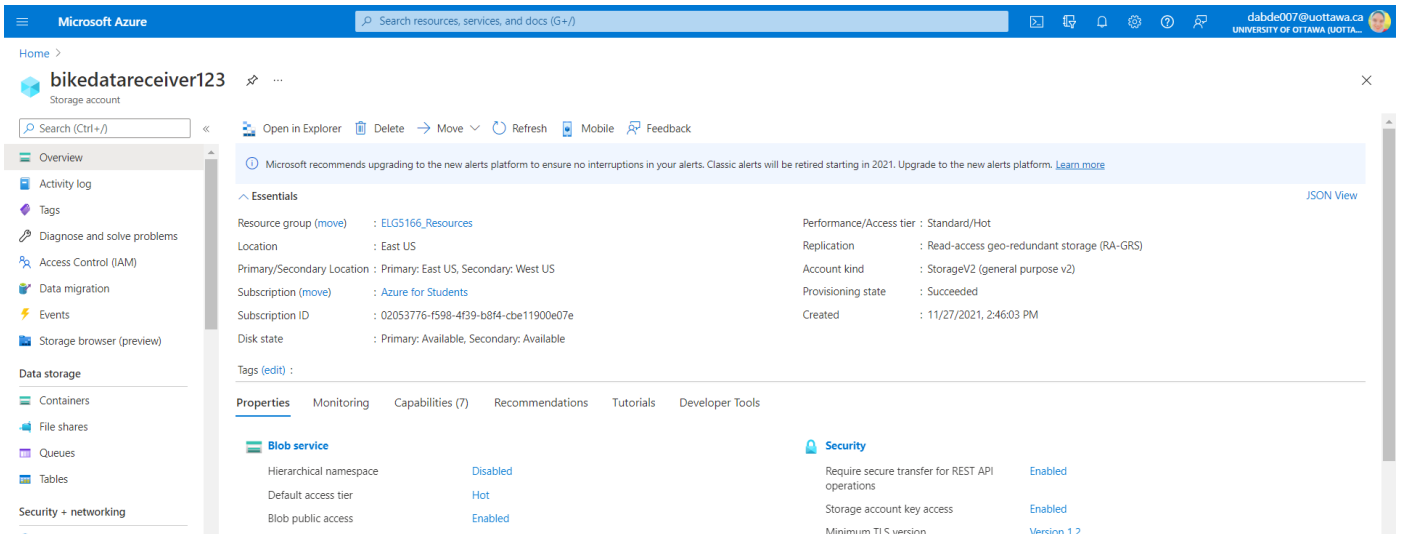
Secondary key
LtpyJ8E44MN4vfywpEAVHhP8+Cae77bewV1PmkP9XE=

Connection string-primary key
Endpoint=sb://assignment3.servicebus.windows.net/SharedA...

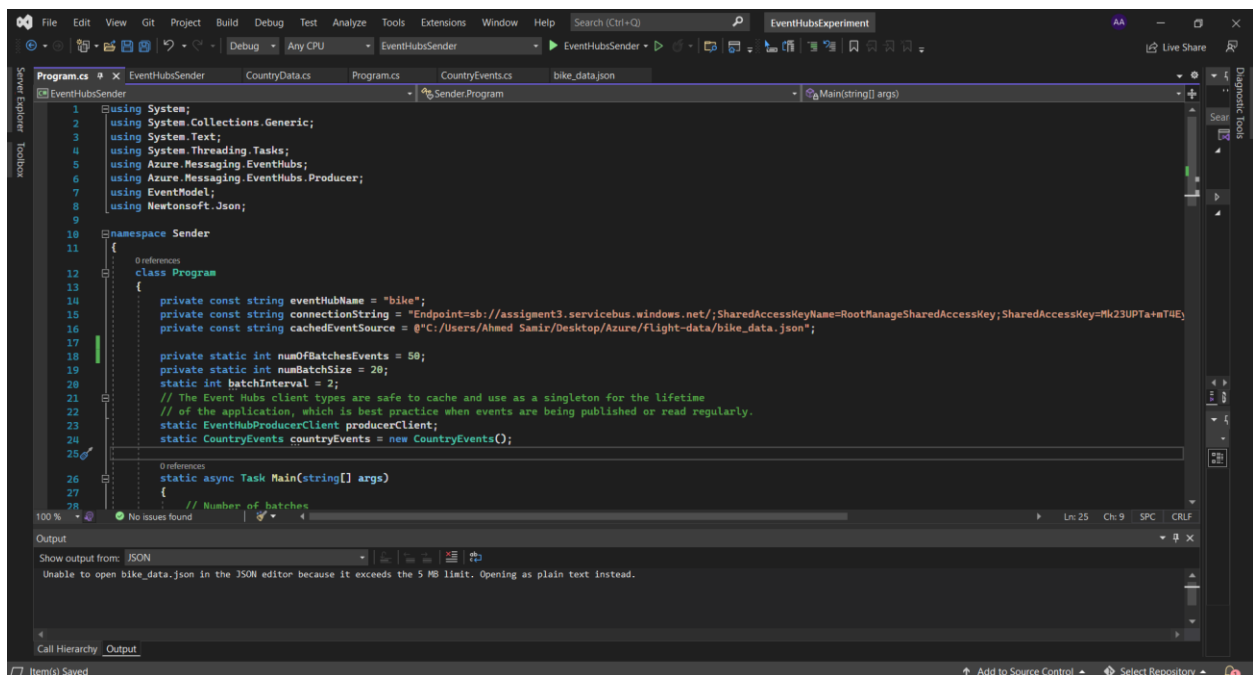
Connection string-secondary key
Endpoint=sb://assignment3.servicebus.windows.net/SharedA...

5. Create storage account

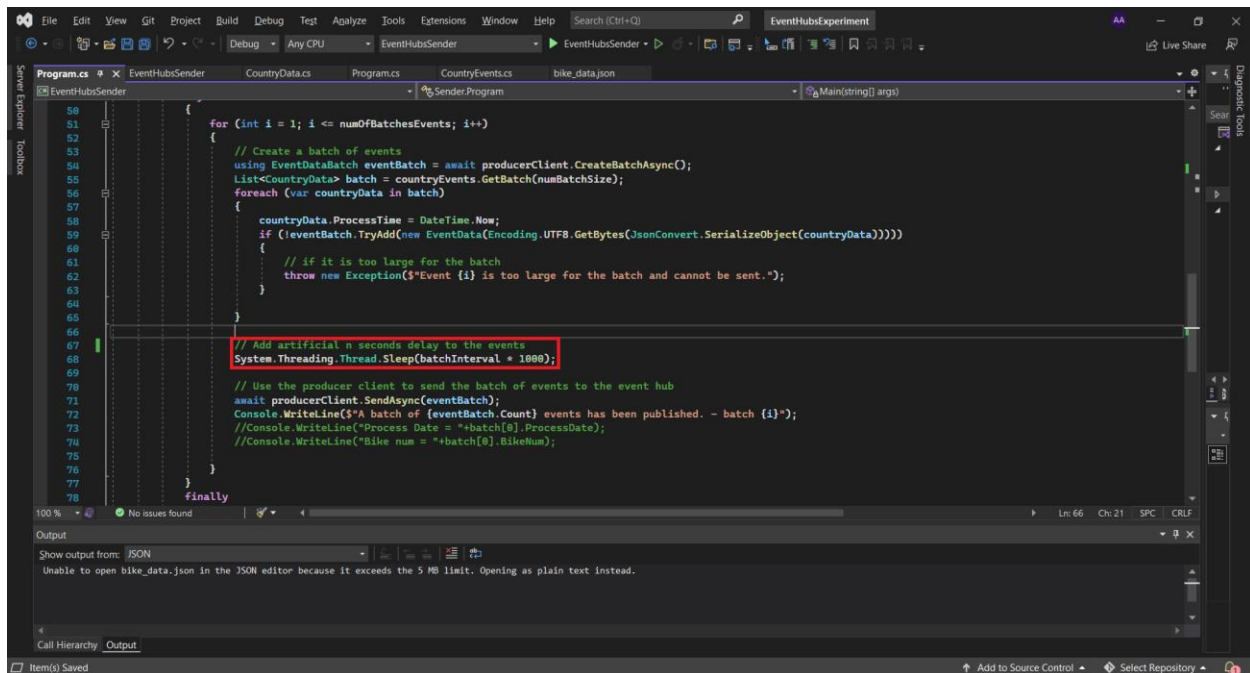
After that we need to change the `eventHubName`, `connectionString` and `cachedEventSource` in the event hub sender to be able to read the data and connect to the Azure.



Also, we need to change the connections on the event hub receiver to be able to receive the data.

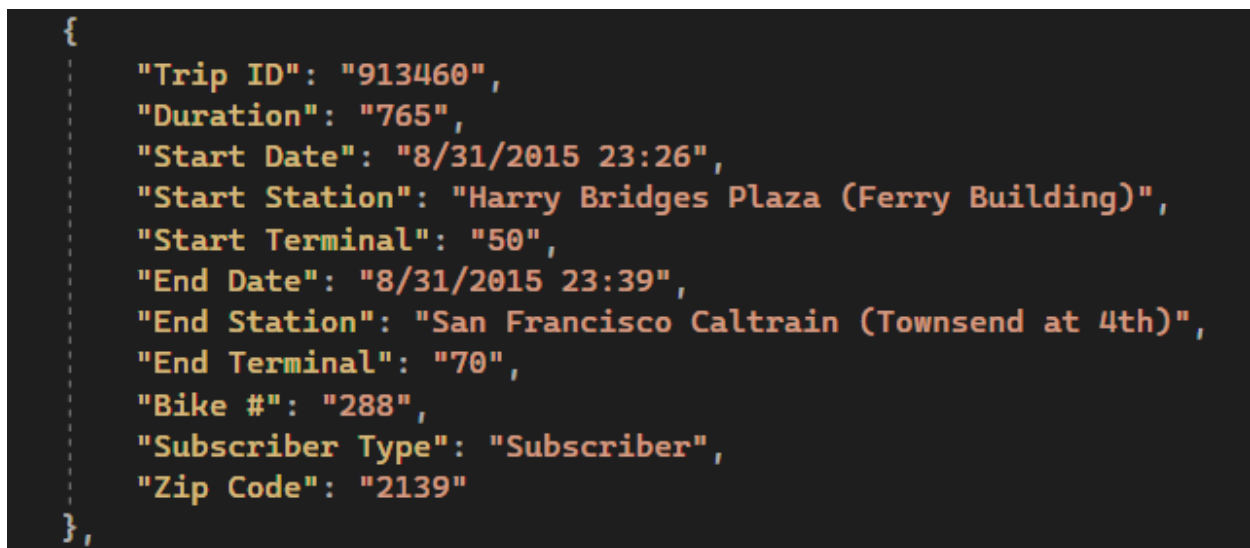


And to change the batch interval we initialized a variable and set its value to 2



```
50 {
51     for (int i = 1; i <= numOfBatchesEvents; i++)
52     {
53         // Create a batch of events
54         using EventDataBatch eventBatch = await producerClient.CreateBatchAsync();
55         List<CountryData> batch = countryEvents.GetBatch(numBatchSize);
56         foreach (var countryData in batch)
57         {
58             countryData.ProcessTime = DateTime.Now;
59             if (!eventBatch.TryAdd(new EventData(Encoding.UTF8.GetBytes(JsonConvert.SerializeObject(countryData)))))
60             {
61                 // if it is too large for the batch
62                 throw new Exception($"Event {i} is too large for the batch and cannot be sent.");
63             }
64         }
65
66         // Add artificial n seconds delay to the events
67         System.Threading.Thread.Sleep(batchInterval * 1000);
68
69         // Use the producer client to send the batch of events to the event hub
70         await producerClient.SendAsync(eventBatch);
71         Console.WriteLine($"A batch of {eventBatch.Count} events has been published. - batch {i}");
72         //Console.WriteLine("Process Date = "+batch[0].ProcessDate);
73         //Console.WriteLine("Bike num = "+batch[0].BikeNum);
74     }
75 }
76 }
77 finally
78 }
```

To send the data we need to look at the Json file

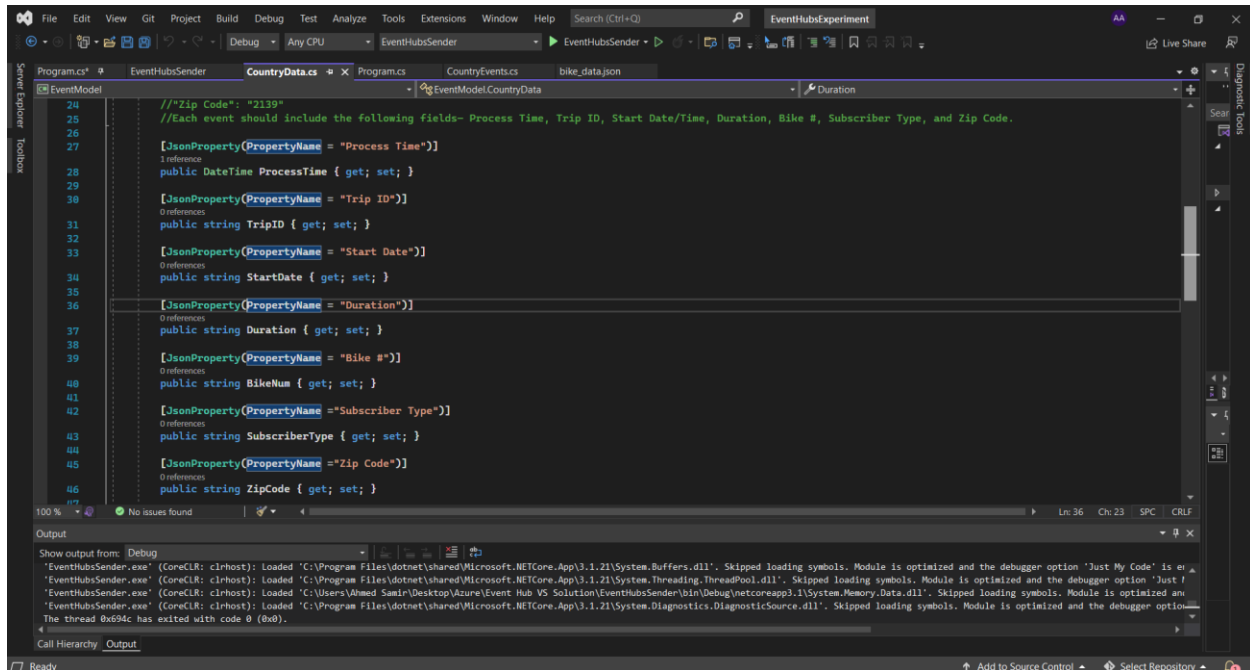


```
{
  "Trip ID": "913460",
  "Duration": "765",
  "Start Date": "8/31/2015 23:26",
  "Start Station": "Harry Bridges Plaza (Ferry Building)",
  "Start Terminal": "50",
  "End Date": "8/31/2015 23:39",
  "End Station": "San Francisco Caltrain (Townsend at 4th)",
  "End Terminal": "70",
  "Bike #": "288",
  "Subscriber Type": "Subscriber",
  "Zip Code": "2139"
},
```

And we can send all the data, but we preferred to send just the data that is required (Process Time, Trip ID, Start Date, Duration, Bike #, Subscriber Type, and Zip Code).

So we changed the class to be able to send these variables and we have added [JsonProperty(PropertyName = variableName)]

to be able to read and send the columns with spaces in between.



The screenshot shows the Visual Studio IDE with the 'CountryData.cs' file open. The class is defined with several properties, each decorated with the [JsonProperty] attribute to specify the name used in the JSON output. The properties are: ProcessTime (DateTime), TripID (string), StartDate (string), Duration (string), BikeNum (string), SubscriberType (string), and ZipCode (string). The output window at the bottom shows the debug output for the EventHubsSender.exe process, indicating that the application is running successfully.

```
//Zip Code: "2139"
//Each event should include the following fields- Process Time, Trip ID, Start Date/Time, Duration, Bike #, Subscriber Type, and Zip Code.

[JsonProperty(PropertyName = "Process Time")]
public DateTime ProcessTime { get; set; }

[JsonProperty(PropertyName = "Trip ID")]
public string TripID { get; set; }

[JsonProperty(PropertyName = "Start Date")]
public string StartDate { get; set; }

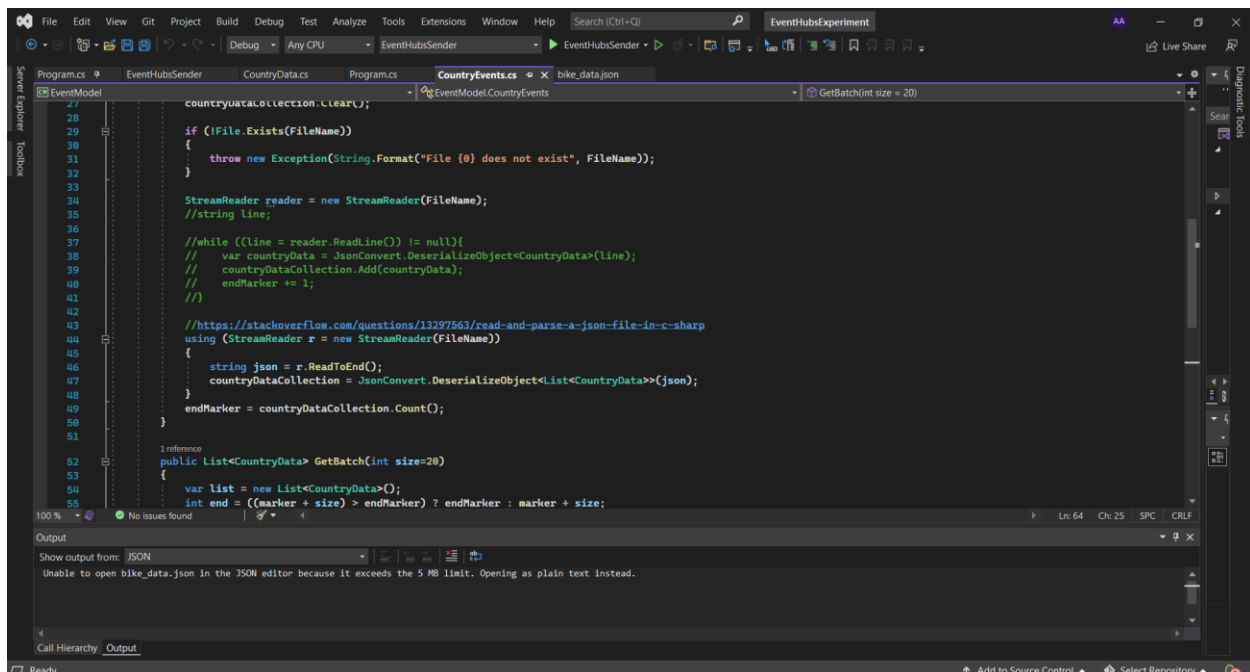
[JsonProperty(PropertyName = "Duration")]
public string Duration { get; set; }

[JsonProperty(PropertyName = "Bike #")]
public string BikeNum { get; set; }

[JsonProperty(PropertyName = "Subscriber Type")]
public string SubscriberType { get; set; }

[JsonProperty(PropertyName = "Zip Code")]
public string ZipCode { get; set; }
```

The initial code is reading the data from the Json file line by line but in our case we can read the entire file at once.



The screenshot shows the Visual Studio IDE with the 'CountryEvents.cs' file open. The class contains a method 'GetBatch' that reads the entire JSON file at once using 'JsonConvert.DeserializeObject'. The output window at the bottom shows the debug output for the EventHubsSender.exe process, indicating that the application is running successfully.

```
countryDataCollection.Clear();

if (!File.Exists(fileName))
{
    throw new Exception(String.Format("File {0} does not exist", fileName));
}

StreamReader reader = new StreamReader(fileName);
//string line;

//while ((line = reader.ReadLine()) != null){
//    var countryData = JsonConvert.DeserializeObject<CountryData>(line);
//    countryDataCollection.Add(countryData);
//    endMarker += 1;
//}

//https://stackoverflow.com/questions/13297563/read-and-parse-a-json-file-in-c-sharp
using (StreamReader r = new StreamReader(fileName))
{
    string json = r.ReadToEnd();
    countryDataCollection = JsonConvert.DeserializeObject<List<CountryData>>(json);
    endMarker = countryDataCollection.Count();
}

1 reference
public List<CountryData> GetBatch(int size=20)
{
    var list = new List<CountryData>();
    int end = ((marker + size) > endMarker) ? endMarker : marker + size;
```

Finally, we can send the batches to the Azure and as we can see it sent 50 batches 1000 record and 50 (20*50)

```
Microsoft Visual Studio Debug Console
Marker = 740
A batch of 20 events has been published. - batch 37
Marker = 760
A batch of 20 events has been published. - batch 38
Marker = 780
A batch of 20 events has been published. - batch 39
Marker = 800
A batch of 20 events has been published. - batch 40
Marker = 820
A batch of 20 events has been published. - batch 41
Marker = 840
A batch of 20 events has been published. - batch 42
Marker = 860
A batch of 20 events has been published. - batch 43
Marker = 880
A batch of 20 events has been published. - batch 44
Marker = 900
A batch of 20 events has been published. - batch 45
Marker = 920
A batch of 20 events has been published. - batch 46
Marker = 940
A batch of 20 events has been published. - batch 47
Marker = 960
A batch of 20 events has been published. - batch 48
Marker = 980
A batch of 20 events has been published. - batch 49
Marker = 1000
A batch of 20 events has been published. - batch 50
Event publishing complete
```

- b) We went to the event hub's namespace and show that the entire message set was received. Provide one or more screenshots.

This screen represents a sample of the receive data

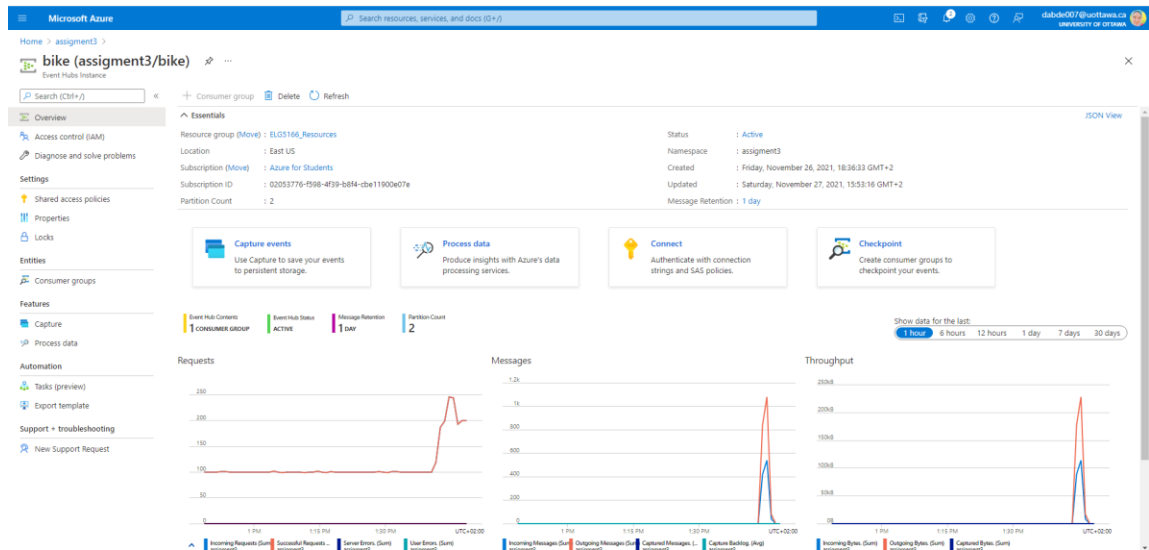
Input preview Test results

Showing events from 'bike'. This list of events might not be complete.


View in JSON Table Raw Refresh Download sample data

Process Time	Trip ID	Start Date	Duration	Bike #	Subscriber Type	Zip Code	EventProcessedUtcTime	PartitionId	EventEnqueuedUtcTime
"2021-11-28T11:43:03.4..."	"913267"	"8/31/2015 19:04"	"780"	"521"	"Subscriber"	"94403"	"2021-11-28T12:04:28.2..."	0	"2021-11-28T11:43:06.6..."
"2021-11-28T11:43:03.4..."	"913266"	"8/31/2015 19:04"	"229"	"41"	"Subscriber"	"94107"	"2021-11-28T12:04:28.2..."	0	"2021-11-28T11:43:06.6..."
"2021-11-28T11:43:01.2..."	"913287"	"8/31/2015 19:17"	"847"	"462"	"Subscriber"	"94107"	"2021-11-28T12:04:28.2..."	1	"2021-11-28T11:43:04.4..."
"2021-11-28T11:43:01.2..."	"913288"	"8/31/2015 19:18"	"199"	"672"	"Subscriber"	"94107"	"2021-11-28T12:04:28.2..."	1	"2021-11-28T11:43:04.4..."
"2021-11-28T11:42:59.0..."	"913293"	"8/31/2015 19:19"	"354"	"502"	"Subscriber"	"94133"	"2021-11-28T12:04:28.2..."	0	"2021-11-28T11:43:02.2..."
"2021-11-28T11:42:59.0..."	"913295"	"8/31/2015 19:19"	"357"	"388"	"Subscriber"	"94109"	"2021-11-28T12:04:28.2..."	0	"2021-11-28T11:43:02.2..."
"2021-11-28T11:42:59.0..."	"913296"	"8/31/2015 19:20"	"765"	"564"	"Subscriber"	"94116"	"2021-11-28T12:04:28.2..."	0	"2021-11-28T11:43:02.2..."
"2021-11-28T11:42:59.0..."	"913298"	"8/31/2015 19:23"	"434"	"505"	"Subscriber"	"94705"	"2021-11-28T12:04:28.2..."	0	"2021-11-28T11:43:02.2..."
"2021-11-28T11:42:59.0..."	"913299"	"8/31/2015 19:23"	"1074"	"434"	"Customer"	"31"	"2021-11-28T12:04:27.2..."	0	"2021-11-28T11:43:02.2..."
"2021-11-28T11:42:59.0..."	"913301"	"8/31/2015 19:25"	"1049"	"608"	"Customer"	"90038"	"2021-11-28T12:04:27.2..."	0	"2021-11-28T11:43:02.2..."

From the event hub's namespace, we can see that the entire message set was received

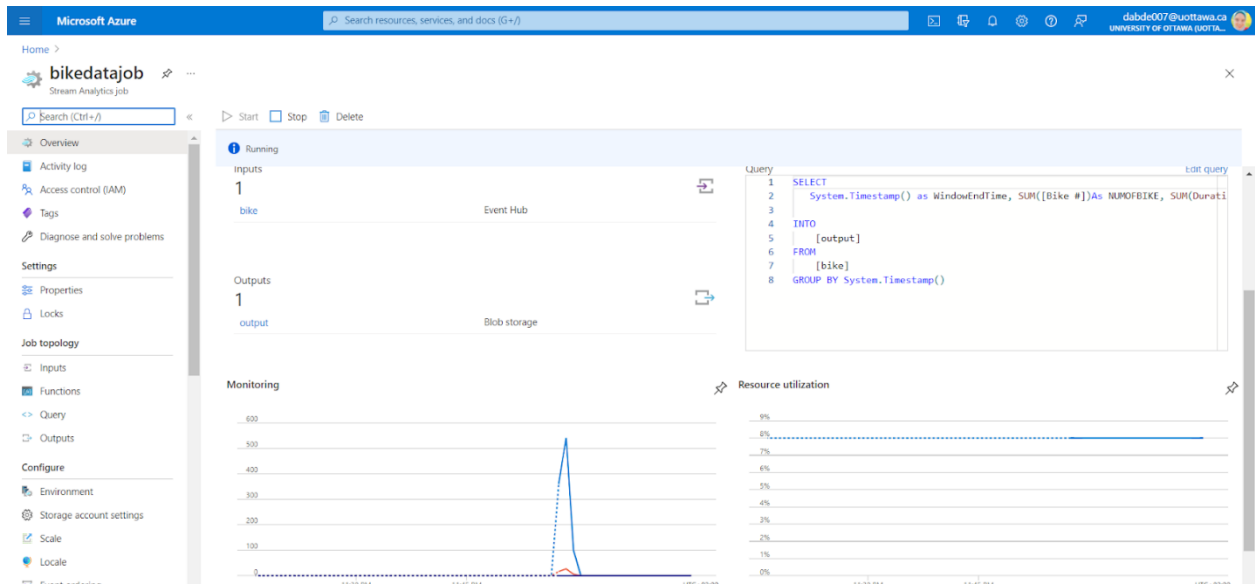


1000 message as 50 events



- c) Then we Created an Azure Stream Analytics job that uses either a Storage account or SQL Database to store the summary of the event received. Our summary should include the following fields – WindowEnd, Total Bikes, Total Duration for each Batch received. Download this dataset and include both a screenshot and the data as part of your submission.

Then, we created stream analytics job



calculate Total Bikes, Total Duration and WindowEnd for each Batch received

Microsoft Azure

Search resources, services, and docs (G+/I)

Home > bikedatajob >

Query

bikedatajob

Query language docs Open in Visual Studio UserVoice

Query can't be edited while a job is running. You can stop the job to edit the query.

Query editor

```
1 SELECT
2   System.Timestamp() as WindowEndTime, SUM([Bike #])As NUMOFBIKE, SUM(Duration) As TotalDuration
3
4 INTO
5   [output]
6 FROM
7   [bike]
8 GROUP BY System.Timestamp()
```

Input preview Test results

Showing events from 'bike'. This list of events might not be complete. Select a specific time range to show all events during that period.

View in JSON Table Raw Refresh Select time range Upload sample input Download sample data

Process Time	Trip ID	Start Date	Duration	Bike #	Subscriber Type	Zip Code	EventProcessedUtcTime	PartitionId	Event
"2021-11-27T23:15:06..."	"913343"	"8/31/2015 19:51"	"426"	"457"	"Subscriber"	"94107"	"2021-11-27T23:20:36..."	1	
"2021-11-27T23:15:06..."	"913344"	"8/31/2015 19:55"	"326"	"672"	"Subscriber"	"94111"	"2021-11-27T23:20:36..."	1	
"2021-11-27T23:15:04..."	"913383"	"8/31/2015 20:20"	"1435"	"604"	"Customer"	"98103"	"2021-11-27T23:20:36..."	0	
"2021-11-27T23:15:04..."	"913382"	"8/31/2015 20:20"	"645"	"270"	"Subscriber"	"94945"	"2021-11-27T23:20:36..."	0	

Data Screenshots:

```
{
  "WindowEndTime": "2021-11-27T21:56:18.2140000Z",
  "NUMOFBIKE": 8294.0,
  "TotalDuration": 10862.0
}
{
  "WindowEndTime": "2021-11-27T21:56:20.5690000Z",
  "NUMOFBIKE": 8520.0,
  "TotalDuration": 9763.0
}
{
  "WindowEndTime": "2021-11-27T21:56:22.8090000Z",
  "NUMOFBIKE": 8536.0,
  "TotalDuration": 11668.0
}
{
  "WindowEndTime": "2021-11-27T21:56:25.1480000Z",
  "NUMOFBIKE": 8090.0,
  "TotalDuration": 9477.0
}
{
  "WindowEndTime": "2021-11-27T21:56:27.3400000Z",
  "NUMOFBIKE": 10100.0,
  "TotalDuration": 9220.0
}
{
  "WindowEndTime": "2021-11-27T21:56:29.5400000Z",
  "NUMOFBIKE": 8968.0,
  "TotalDuration": 11335.0
}
{
  "WindowEndTime": "2021-11-27T21:56:31.7310000Z",
  "NUMOFBIKE": 8845.0,
  "TotalDuration": 8831.0
}
{
  "WindowEndTime": "2021-11-27T21:56:33.9150000Z",
  "NUMOFBIKE": 8616.0,
  "TotalDuration": 10821.0
}
{
  "WindowEndTime": "2021-11-27T21:56:36.1060000Z",
  "NUMOFBIKE": 9729.0,
  "TotalDuration": 22219.0
}
{
  "WindowEndTime": "2021-11-27T21:56:38.2910000Z",
  "NUMOFBIKE": 8887.0,
  "TotalDuration": 10861.0
}
{
  "WindowEndTime": "2021-11-27T21:56:40.4660000Z",
  "NUMOFBIKE": 9461.0,
  "TotalDuration": 11437.0
}
{
  "WindowEndTime": "2021-11-27T21:56:42.6510000Z",
  "NUMOFBIKE": 7326.0,
  "TotalDuration": 9858.0
}
{
  "WindowEndTime": "2021-11-27T21:56:44.8260000Z",
  "NUMOFBIKE": 8878.0,
  "TotalDuration": 11478.0
}
{
  "WindowEndTime": "2021-11-27T21:56:47.2610000Z",
  "NUMOFBIKE": 7311.0,
  "TotalDuration": 10866.0
}
{
  "WindowEndTime": "2021-11-27T21:56:49.4370000Z",
  "NUMOFBIKE": 9042.0,
  "TotalDuration": 10990.0
}
{
  "WindowEndTime": "2021-11-27T21:56:51.6360000Z",
  "NUMOFBIKE": 8027.0,
  "TotalDuration": 9910.0
}
{
  "WindowEndTime": "2021-11-27T21:56:53.8280000Z",
  "NUMOFBIKE": 8734.0,
  "TotalDuration": 11262.0
}
{
  "WindowEndTime": "2021-11-27T21:56:55.9950000Z",
  "NUMOFBIKE": 9529.0,
  "TotalDuration": 10518.0
}
{
  "WindowEndTime": "2021-11-27T21:56:58.1880000Z",
  "NUMOFBIKE": 8080.0,
  "TotalDuration": 9852.0
}
{
  "WindowEndTime": "2021-11-27T21:57:00.3620000Z",
  "NUMOFBIKE": 9409.0,
  "TotalDuration": 11102.0
}
{
  "WindowEndTime": "2021-11-27T21:57:02.5480000Z",
  "NUMOFBIKE": 7880.0,
  "TotalDuration": 10662.0
}
{
  "WindowEndTime": "2021-11-27T21:57:04.7840000Z",
  "NUMOFBIKE": 7721.0,
  "TotalDuration": 11399.0
}
{
  "WindowEndTime": "2021-11-27T21:57:06.9700000Z",
  "NUMOFBIKE": 7830.0,
  "TotalDuration": 10598.0
}
{
  "WindowEndTime": "2021-11-27T21:57:09.1590000Z",
  "NUMOFBIKE": 8498.0,
  "TotalDuration": 11332.0
}
{
  "WindowEndTime": "2021-11-27T21:57:11.3460000Z",
  "NUMOFBIKE": 7687.0,
  "TotalDuration": 13254.0
}
{
  "WindowEndTime": "2021-11-27T21:57:13.5180000Z",
  "NUMOFBIKE": 8614.0,
  "TotalDuration": 17192.0
}
{
  "WindowEndTime": "2021-11-27T21:57:15.7050000Z",
  "NUMOFBIKE": 7785.0,
  "TotalDuration": 10032.0
}
{
  "WindowEndTime": "2021-11-27T21:57:17.8940000Z",
  "NUMOFBIKE": 8985.0,
  "TotalDuration": 11039.0
}
{
  "WindowEndTime": "2021-11-27T21:57:20.0810000Z",
  "NUMOFBIKE": 9339.0,
  "TotalDuration": 9578.0
}
{
  "WindowEndTime": "2021-11-27T21:57:22.2840000Z",
  "NUMOFBIKE": 7240.0,
  "TotalDuration": 12372.0
}
{
  "WindowEndTime": "2021-11-27T21:57:24.4570000Z",
  "NUMOFBIKE": 9077.0,
  "TotalDuration": 19120.0
}
{
  "WindowEndTime": "2021-11-27T21:57:26.7380000Z",
  "NUMOFBIKE": 8323.0,
  "TotalDuration": 12728.0
}
{
  "WindowEndTime": "2021-11-27T21:57:29.0190000Z",
  "NUMOFBIKE": 8806.0,
  "TotalDuration": 75957.0
}
{
  "WindowEndTime": "2021-11-27T21:57:31.3310000Z",
  "NUMOFBIKE": 8888.0,
  "TotalDuration": 60530.0
}
{
  "WindowEndTime": "2021-11-27T21:57:33.5200000Z",
  "NUMOFBIKE": 8308.0,
  "TotalDuration": 9954.0
}
{
  "WindowEndTime": "2021-11-27T21:57:35.7070000Z",
  "NUMOFBIKE": 7586.0,
  "TotalDuration": 11852.0
}
{
  "WindowEndTime": "2021-11-27T21:57:37.9570000Z",
  "NUMOFBIKE": 7638.0,
  "TotalDuration": 13610.0
}
{
  "WindowEndTime": "2021-11-27T21:57:40.1290000Z",
  "NUMOFBIKE": 7909.0,
  "TotalDuration": 40151.0
}
{
  "WindowEndTime": "2021-11-27T21:57:42.3170000Z",
  "NUMOFBIKE": 9013.0,
  "TotalDuration": 8945.0
}
{
  "WindowEndTime": "2021-11-27T21:57:44.4890000Z",
  "NUMOFBIKE": 8551.0,
  "TotalDuration": 9157.0
}
{
  "WindowEndTime": "2021-11-27T21:57:46.6770000Z",
  "NUMOFBIKE": 9371.0,
  "TotalDuration": 10497.0
}
{
  "WindowEndTime": "2021-11-27T21:57:48.8800000Z",
  "NUMOFBIKE": 7349.0,
  "TotalDuration": 8131.0
}
{
  "WindowEndTime": "2021-11-27T21:57:51.0670000Z",
  "NUMOFBIKE": 8612.0,
  "TotalDuration": 13008.0
}
{
  "WindowEndTime": "2021-11-27T21:57:53.2400000Z",
  "NUMOFBIKE": 9180.0,
  "TotalDuration": 10914.0
}
{
  "WindowEndTime": "2021-11-27T21:57:55.4280000Z",
  "NUMOFBIKE": 9384.0,
  "TotalDuration": 11241.0
}
{
  "WindowEndTime": "2021-11-27T21:57:57.6190000Z",
  "NUMOFBIKE": 7845.0,
  "TotalDuration": 11318.0
}
{
  "WindowEndTime": "2021-11-27T21:57:59.8030000Z",
  "NUMOFBIKE": 7269.0,
  "TotalDuration": 9608.0
}
{
  "WindowEndTime": "2021-11-27T21:58:01.9940000Z",
  "NUMOFBIKE": 8511.0,
  "TotalDuration": 11939.0
}
{
  "WindowEndTime": "2021-11-27T21:58:04.1780000Z",
  "NUMOFBIKE": 7629.0,
  "TotalDuration": 11025.0
}
{
  "WindowEndTime": "2021-11-27T21:58:06.3700000Z",
  "NUMOFBIKE": 8439.0,
  "TotalDuration": 11456.0
}
```

We can see that we have 50 record each record represents the number of bikes and the total duration per batch.

2. Azure Data Lake Analytics

2.1. Create a Data Lake Storage Gen2 account

To create ADL (Azure Data Lake), It is required to create Data Lake Storage in Figure-1

Microsoft Azure portal showing the 'New Data Lake Storage Gen1' creation page. The page includes a sidebar with 'Data Lake Storage Gen1' and a main form with fields for Subscription, Resource group, Location, and Name. A warning banner at the top states: 'Azure Data Lake Storage Gen1 will be retired on February 29, 2024. We recommend that you migrate your Azure Data Lake Storage Gen1 account to Azure Data Lake Storage Gen2.' The form is partially filled with 'Azure for Students' for Subscription, 'ELG5166_Resources' for Resource group, 'East US 2' for Location, and 'assignment3' for Name. The URL in the browser is 'portal.azure.com/#create/MicrosoftAzureDataLakeStorage'.

Figure 1- Create a Data Lake Storage Gen2 account

2.2. Create Data Analytics Account

Microsoft Azure portal showing the 'New Data Lake Analytics account' creation page. The page includes a sidebar with 'Data Lake Analytics' and a main form with fields for Subscription, Resource group, Location, Name, Existing storage subscription, and Pricing package. The form is partially filled with 'Azure for Students' for Subscription, 'ELG5166_Resources' for Resource group, 'East US 2' for Location, 'assignment3' for Name, 'Azure for Students' for Existing storage subscription, and 'Pay-as-you-go' for Pricing package. The URL in the browser is 'portal.azure.com/#create/MicrosoftAzureDataLakeAnalytics'.

Figure 2- Set up New Data Lake Analytics Account

2.3. Microsoft Azure Storage Explorer

Moving our data samples to the ADLS requires us first to create the Input and Output folders to upload the rest of the content, So Microsoft Azure Store Explorer was used.

Here in Figure 3, We have finished our connection with Microsoft Azure Storage Explorer with our local data by uploading Input and Output folders and by uploading all DLL files to connect Microsoft Visual Studio with Azure Data Lake Analytics to have live updates over cloud from our local machine.

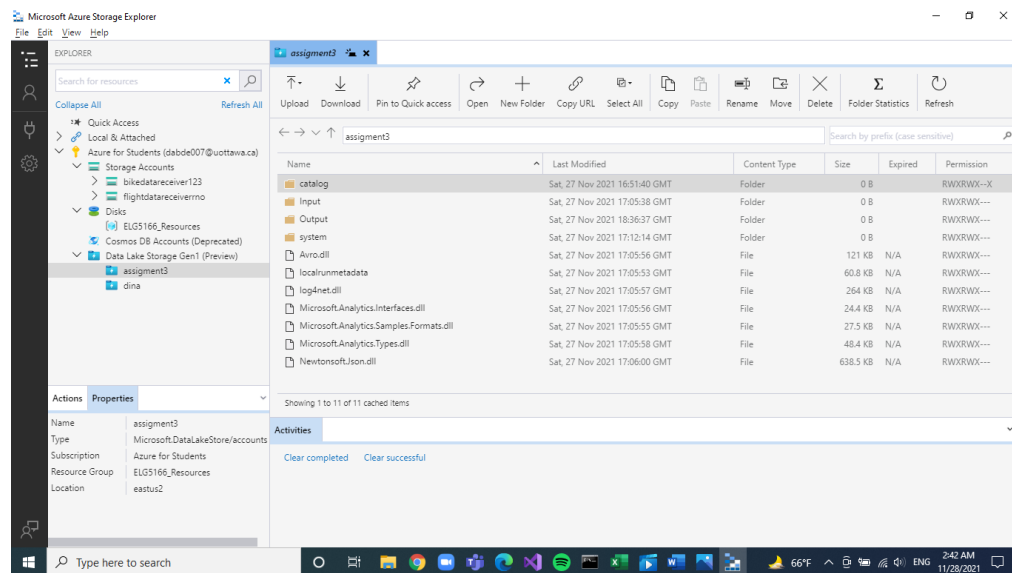


Figure 3 - Microsoft Azure Storage Explorer Finished Setup

2.4. U-SQL Queries

The tables in Tabel 1 and Table 2 represents a bikes rental data, In table 1, It represents bike trips the was done at specific time from starting terminal to end terminal having specific amount of bikes and the type of customers that had that trip that are living at specific zip code as their location.

In table 2, It describes the terminals information.

In this Assignment, We want to derive some meaningful information from this data.

Trip ID	Durati on	Start Date	Start Terminal	End Date	End Terminal	Bike #	Subscriber Type	Zip Code
913460	765	8/31/2015 23:26	50	8/31/2015 23:39	70	288	Subscriber	2139
913459	1036	8/31/2015 23:11	31	8/31/2015 23:28	27	35	Subscriber	95032
913455	307	8/31/2015 23:13	47	8/31/2015 23:18	64	468	Subscriber	94107
913454	409	8/31/2015 23:10	10	8/31/2015 23:17	8	68	Subscriber	95113
913453	789	8/31/2015 23:09	51	8/31/2015 23:22	60	487	Customer	9069

Table 1- Bikes Trips

```
@Trips = EXTRACT
[Trip ID] string,
[Duration] int,
[Start Date] DateTime,
[Start Terminal] int,
[End Date] DateTime,
[End Terminal] int,
[Bike #] int,
[Subscriber Type] string,
[Zip Code] string
FROM "/Input/201508_trip_data.csv"
USING Extractors.Csv(skipFirstNRows:1);
```

Figure 4 - Trips Table in U-SQL

station_id	name	lat	long	dockcount	landmark	installation
2	San Jose Diridon Caltrain Station	37.32973	-121.902	27	San Jose	8/6/2013
3	San Jose Civic Center	37.3307	-121.889	15	San Jose	8/5/2013
4	Santa Clara at Almaden	37.33399	-121.895	11	San Jose	8/6/2013
5	Adobe on Almaden	37.33142	-121.893	19	San Jose	8/5/2013

Table 2 - Terminal Data

```

@Station = EXTRACT
[station_id] int,
[name] string,
[lat] float,
[long] float,
[dockcount] int,
[landmark] string,
[installation] DateTime
FROM "/Input/201508_station_data.csv"
USING Extractors.Csv(skipFirstNRows:1);

```

Figure 5 - Stations Table in U-SQL

Queries

1. Top 20 zip codes where most bikes were rented from

Assumptions: we have assumed [Bike #] is the number of bikes for one trip because it was mentioned “#” right after the “Bike”.

From **Table 1**, the sum over [Bike #] was calculated then group by the [zip code] followed by Order by the sum of bikes descending, and finally fetching the first 20 rows.

Result: In table 4, we have the highest 20 zip codes that have bike number.

```
@query_1 =  
SELECT SUM([Bike #]) AS Bikes_Number,  
       [Zip Code]  
FROM @Trips  
GROUP BY [Zip Code]  
ORDER BY Bikes_Number DESC  
FETCH 20 ROWS;  
  
OUTPUT @query_1  
TO "/Output/trip_data_1.csv"  
USING Outputters.Csv(outputHeader : true);
```

Figure 6- Top 20 zip codes query

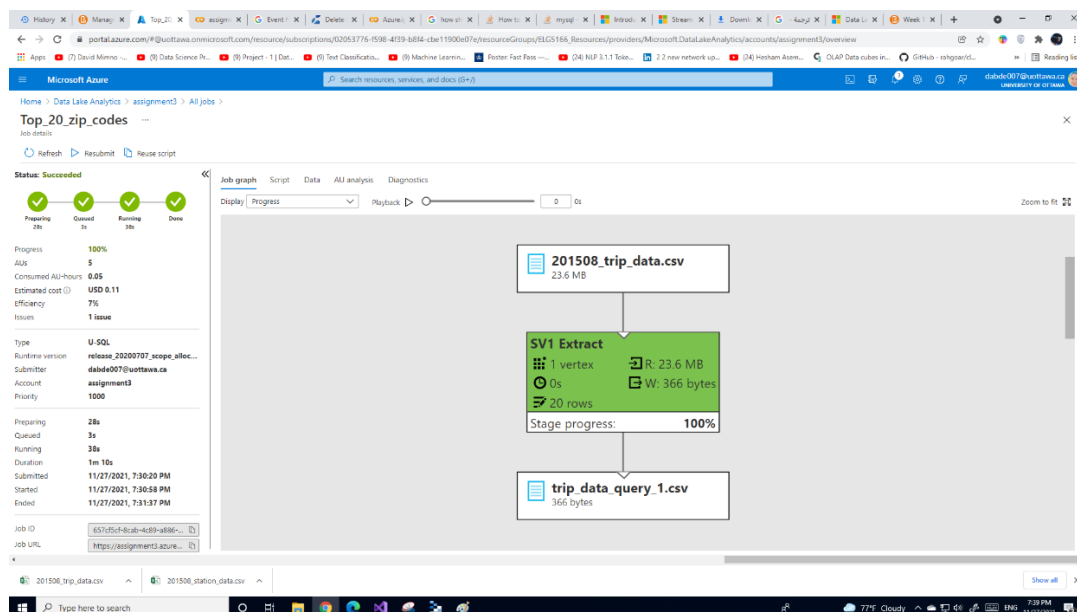


Figure 7 - Query 1 graph

Zip Code	Bikes_Number
94107	20228868
94105	8788689
94133	7077405
94103	6353674
94111	4782022

Table 3 - Top 20 zip codes results

2. Daily duration aggregate across the rental subscriber types.

From **Table 1**, The sum over the single day trip durations was calculated separately upon each type by using CASE statement for each subscriber type, then grouping by the day duration was performed

Result: In Table 4, The duration of all trips of each day for both subscriber types.

```
@query_2 =
SELECT [Start Date].ToString("MM/dd/yyyy") AS Date ,
SUM([Subscriber Type] == "Subscriber" ? [Duration] : 0)AS Total_Duration_Subscribers,
SUM([Subscriber Type] == "Customer" ? [Duration] : 0 )AS Total_Duration_Customers
FROM @Trips
GROUP BY [Start Date].ToString("MM/dd/yyyy");

OUTPUT @query_2
TO "C:/Users/elsha/Downloads/assignment 3/Output/trip_data_query_2.csv"
USING Outputters.Csv(outputHeader : true);
```

Figure 8 - Daily duration aggregation Query

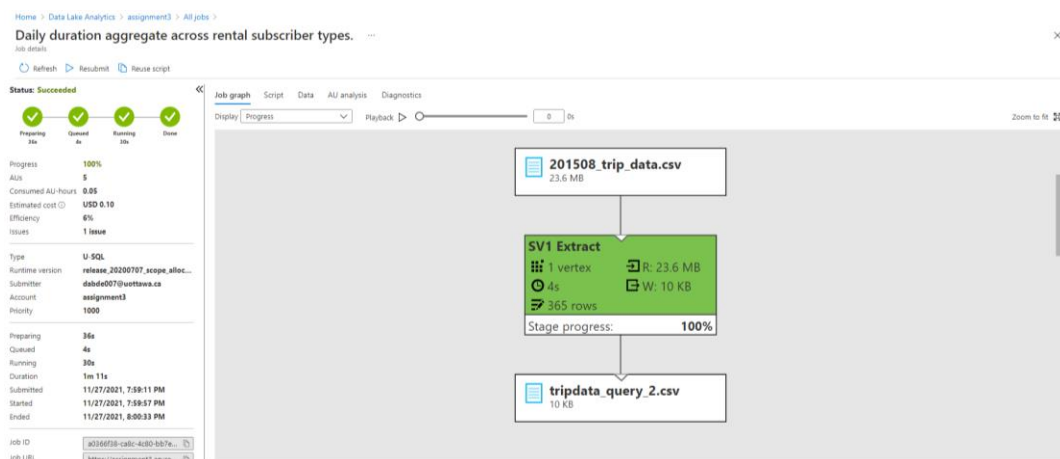


Figure 9 - Daily Duration Aggregation graph

Date	Total_Duration_Subscribers	Total_Duration_Customers
1/1/2015	37943	344997
1/2/2015	154422	772045
1/3/2015	80711	460716
1/4/2015	68021	489093
1/5/2015	594816	159695
1/6/2015	625803	187769

Table 4 - Daily Duration aggregation results

3. The top busiest terminals for bike pickup

Assumption: we have assumed that bike pickup means that the terminals that the bikes started their trip from, hence the start terminal was used in the query, And the busiest start terminal would have the most occurrences in the data.

From **Table 1**, we have grouped by the [start terminal] and selecting the count of all terminal occurrences to have the frequency of all start terminals and selected the start terminal number.

Then we have performed an “INNER JOIN” with **Table 2** with the [station_id] to get the names of the terminals from **Table 2**, and lastly ordering by the frequency of terminals [busy_rate] descending and then fetching the top 20 records.

```
@query_3 =
    SELECT s.station_id AS Station_ID,
           s.name AS Station_name,
           t.busy_rate
    FROM
        ( SELECT [Start Terminal] AS start_terminal,
                  COUNT([Start Terminal]) AS busy_rate
          FROM @Trips
          GROUP BY [Start Terminal] ) AS t
    INNER JOIN
        @Station AS s
    ON s.station_id == t.start_terminal
    ORDER BY busy_rate DESC
    FETCH 20 ROWS;

OUTPUT @query_3
TO "/Output/tripdata_query_3.csv"
USING Outputters.Csv(outputHeader : true);
```

Table 5 - top most busy start terminals query

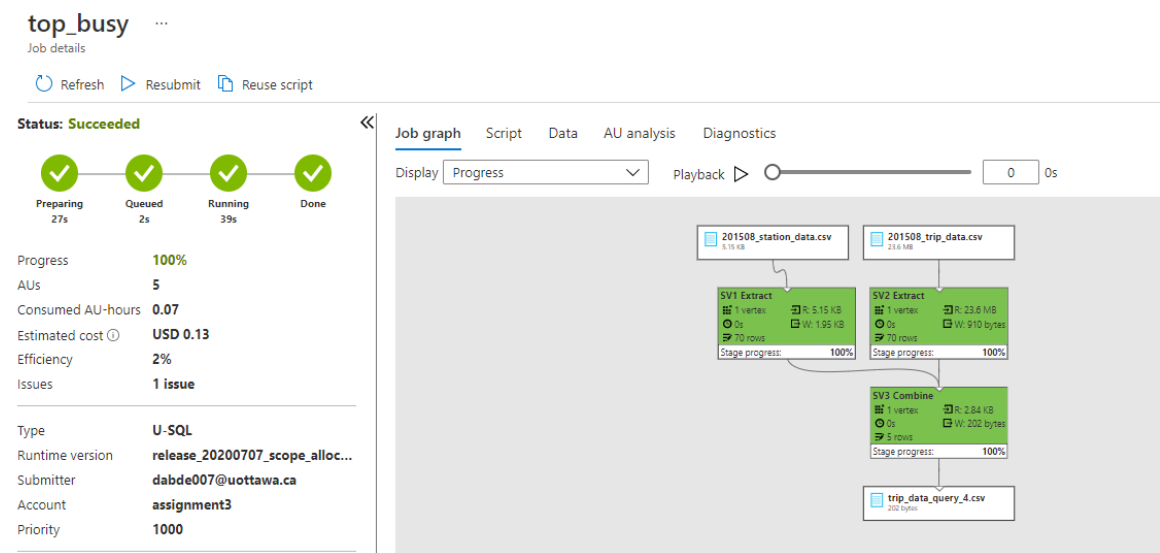


Table 6 - top most busy start terminals graph

Station_ID	Station_name	busy_rate
70	San Francisco Caltrain (Townsend at 4th)	26304
69	San Francisco Caltrain 2 (330 Townsend)	21758
50	Harry Bridges Plaza (Ferry Building)	17255
55	Temporary Transbay Terminal (Howard at Beale)	14436
60	Embarcadero at Sansome	14158

Table 7 - top most busy start terminals result

4. Which 5 terminal has the least drop-offs?

Assumption: we have assumed that bike drop-offs means that the terminals that the bikes ended their trip on, hence the [end terminal] was used in the query, And the least dropoff terminals would have the least occurrences in the dataset.

From **Table 1**, we have grouped by the [end terminal] and selecting the count of all terminal occurrences to have the frequency of all end terminals and selected the end terminal number.

Then we have performed an "INNER JOIN" with **Table 2** with the [station_id] to get the names of the terminals from **Table 2**, and lastly ordering by the frequency of terminals [Dropoff_Rate] ascending and then fetching the top 5 records.

```

@query_4 =
SELECT s.station_id AS Station_ID, s.name AS Station_name,
      t.Dropoff_Rate
FROM
(
    SELECT [End Terminal] AS end_terminal,
           COUNT([End Terminal]) AS Dropoff_Rate
    FROM @Trips
    GROUP BY [End Terminal]
) AS t
INNER JOIN
    @Station AS s
ON s.station_id == t.end_terminal
ORDER BY Dropoff_Rate ASC
FETCH 5 ROWS;

OUTPUT @query_4
TO "/Output/tripdata_query_4.csv"
USING Outputters.Csv(outputHeader : true);

```

Figure 10 - least terminal dropoffs query

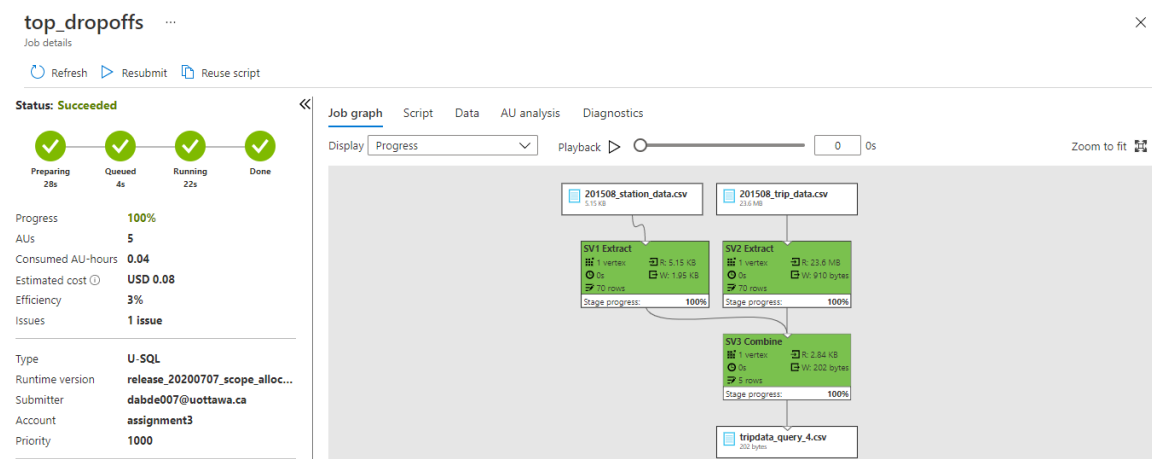


Figure 11 - least terminal dropoffs graph

Station_ID	Station_name	Dropoff_Rate
24	Redwood City Public Library	98
21	Franklin at Maple	100
83	Mezes Park	145
23	San Mateo County Center	187
26	Redwood City Medical Center	230

Table 8 - least terminal dropoffs results

5. What is the monthly summary of bike rentals (format - month/year ex. 06/2020)

Total bikes, total duration (seconds), total trips and total subscriber types was calculated per month.

```
@query_5 =  
SELECT [Start Date].ToString("yyyy/MM") AS Date, COUNT([Trip ID]) AS Trip_Count,  
SUM([Duration]) AS Total_Duration,  
SUM([Bike #]) AS Total_Bikes,  
SUM([Subscriber Type] == "Subscriber" ? 1 : 0) AS Total_Subscribers,  
SUM([Subscriber Type] == "Customer" ? 1 : 0) AS Total_Customers  
FROM @Trips  
GROUP BY [Start Date].ToString("yyyy/MM");  
  
OUTPUT @query_5  
TO "/Output/trip_data_query_5.csv"  
USING Outputters.Csv(outputHeader : true);
```

Figure 12 - summary query

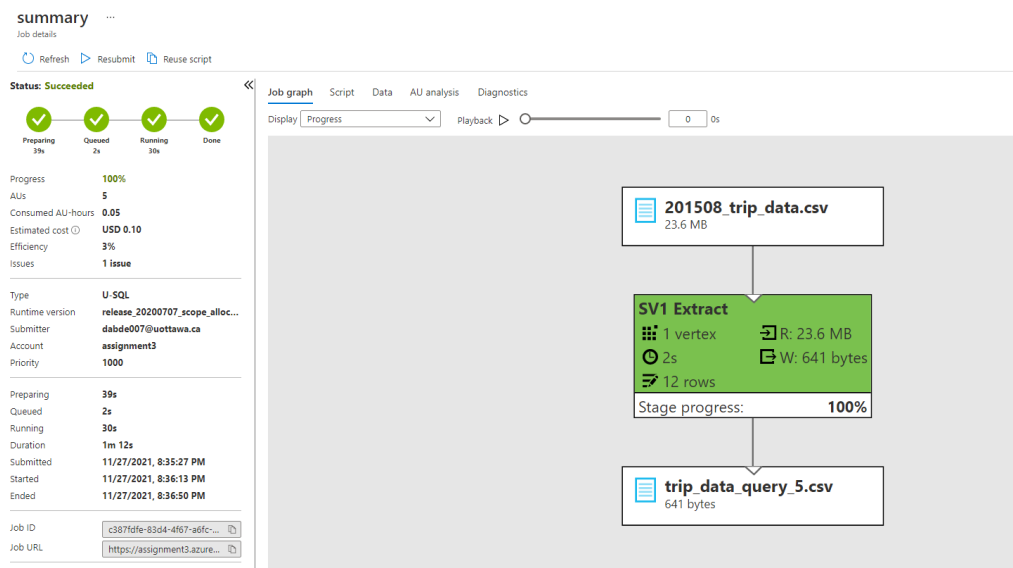


Figure 13 - summary graph

Date	Trip_Count	Total_Duration	Total_Bikes	Total_Subscribers	Total_Customers
2014/09	31682	33160021	13310244	27456	4226
2014/10	34220	33401099	14386280	29833	4387
2014/11	25516	22454934	10752350	22623	2893
2014/12	19677	41131402	8382802	17386	2291
2015/01	27840	25611358	11756577	25068	2772
2015/02	26401	25633016	11139951	23688	2713
2015/03	31626	29892301	13444028	27752	3874
2015/04	31363	28031940	13313721	28038	3325
2015/05	29540	31584633	12574550	25545	3995
2015/06	31907	34481927	13447867	27868	4039
2015/07	32476	33983062	13718375	27652	4824
2015/08	31904	31088866	13564645	27308	4596

Table 9 - Summary result

3. Definitions

1. Please compare briefly, based on at least 3 criteria, the differences in architecture between Apache Spark Structured Streaming and Azure Event Hubs & Stream Analytics?

Event hub and stream analytics work together as end-to-end solution where the event hub read the stream and feed it to the stream analytics to analyze this data. while spark input File source, Socket Source, Rate Source, Kafka Source

In stream analytics User can extend existing functions as UDFs in JavaScript or C# or User-defined aggregates. While in spark it uses Scala Python R Java

The input data from in-stream analytics Avro, JSON or CSV, UTF-8 encoded, while spark is taking any format using custom code

References:

- [X]Microsoft Docs, <https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/data/stream-processing-stream-analytics> , Stream processing with Azure Stream Analytics
- [X] Microsoft Docs, <https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/stream-processing> ,Choose a stream processing technology in Azure
- [X] Spark apache , <https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html> ,Structured Streaming Programming Guide

2. Describe briefly 3 benefits of Azure U-SQL over Apache Spark. Illustrate them briefly with some use cases

Although U-SQL is a combination between spark and SQL and it is not an ANSI SQL, It supports constructs and U-SQL has some advantages over spark:

- The first thing is that U-SQL makes the user write normal SQL queries and can be driven to C#. This helps to extend it to us to user-defined functions. C# is statically typed, and it is faster, and it has can do more than python and this

is useful in the case of streaming. While spark use with Scala or python, while U-SQL can extend with c# and .Net it allows us to:

- Add c# expressions inside SELECT.
- Add user-defined functions, and aggregation,
- Add user-defiend operators like (processors, extractors, outputters, reducers and combiners).
- U-SQL cost per query. but Apache spark costs keep going as long as the cluster is running by U-SQL we can save more money. As if we. have a large data it's highly recommended to get payed on the job you have done only.
- U-SQL could be used on the local machine in visual basic this helps to show the results before going to the cloud. This could be helpful to monitor the behavior of the query before going to the cloud if something went wrong then we can modify it.
- Uniformally process the data from different sources in Azure (like Azure data lake, Azure block storage, SQL DB, SQL Server instances, and ADLS), so it will be better to use U-SQL on the Azure environment.

References:

- Lecture notes
- [X] stack overflow,
<https://stackoverflow.com/questions/35575080/azure-spark-sql-vs-u-sql> , Azure Spark SQL vs U-SQL

3. What are the 5 characteristics of Azure Data Lake that distinguish it from other Distributed Dataset Storage infrastructure such as Hadoop?

1. Developer friendly platform

Managing Hadoop applications is complex. Developers need to write code to each operation that makes it very difficult to work but Azure Data Lake makes this easy by using Azure DevOps, Visual Studio, etc..., so that make Developers use familiar tools to run, debug and trace the code. And see how it runs.

2. Security

Data in the cluster is not always segmented by applications. Hadoop clusters don't support encrypted data at rest. But Azure Data Lake is Based on HDFS but provides security management & integration with other Azure services while abstracting Hadoop commands. So, it is managed and supported by Microsoft and governance controls.

3. Unlimited Storage

Azure Data Lake Storage can hold any amount of data and unlimited number of files. We can store single or multiple files at terabyte and petabyte. Because the files are divided into blocks. these blocks are distributed across data nodes. Which are unlimited. Other cloud storage options have file and account size limitations of only a few terabytes.

4. Real-time Data Processing

Apache Hadoop is for batch processing, which means it takes a huge amount of data in input, process it and produces the result. But Azure Data Lake has Apache Spark/Storm in HDInsight which processes streaming data for the Azure DL Storage.

5. the data into one place

ADLS can gathering all the big data from disparate sources across cloud and on-premises environments into one place. That make the monitoring and managing the stored data easy.

References:

[1]Lecture Notes , Week 8 - Azure Big Data Landscape & Week 4 - Batch Processing of Big Data

[2] dremio, <https://www.dremio.com/data-lake/azure/> -What Is Azure Data Lake Storage (ADLS)?

[3]Dataflair, <https://data-flair.training/blogs/13-limitations-of-hadoop/> -13 Big Limitations of Hadoop & Solution To Hadoop Drawbacks

4. What 2 factors influence the use of Azure AUs for U-SQL query processing. Describe with a simple example.

AU number of nodes to process the job the default AUs number on azure is 5 there is a tool called AU Analyzer that helps to perform optimizations. So, depending on the job we want to make you can choose the number of AUs. For example, if we process a huge amount of data then we need to add more AUs. and if we want to make simple things as creating a table then 1 AU could be enough. The two factors that influence choosing the number of AUs are cost and time. but more AUs can make the overhead due to parallelization

For example, if we need to query a huge amount of data in use this data in some real-time application so it is highly recommended to increase the number of AUs. And if we want to get some insight from the data lake (Time not important) we can choose a fewer number of AUs.

References:

- Lecture notes, recording
- [X] azure, <https://azure.microsoft.com/en-au/blog/get-started-with-u-sql-it-s-easy/> ,Get started with U-SQL: It's easy!